



SEATTLE TRAFFICK DATA ANALYSIS AND RISK PREDICTION

IBM Final Capstone Project

Puja Misra

Introduction

This project is our Final submission to IBM Data Science Professional Certificate course on Coursera. The goal of the project is to detail and use Data Science tool set for Predictive analysis.

We will be working on a real-life problem and demonstrate how Machine Learning can help us predict and process the value by applying the learned skills.

Business Understanding

- **Background:**

According to 2017 WSDOT data, a car accident occurs every 4 minutes and a person dies due to a car crash every 20 hours. Fatal crashes went from 508 in 2016 to 525 in 2017, resulting in the death of 555 people. This number has stayed relatively steady for the past decade. According to 2017 WSDOT data, a car accident occurs every 4 minutes and a person dies due to a car crash every 20 hours. Fatal crashes went from 508 in 2016 to 525 in 2017, resulting in the death of 555 people. This number has stayed relatively steady for the past decade.

- **Problem Statement:**

As we see in the background statement above, the numbers of fatal crashes are having an upward trend or has been steady for past decade for Seattle as per WSDOT .



- **Objective of the Project:**

- ✚ The purpose of the project is to gather the data and determine what causes the accident and the attributes that leads to the severity.
- ✚ Through data visualization and machine learning algorithm we will be analyzing a significant range of attributes, including weather conditions, road condition, speeding, special events, roadworks, traffic jams among others and we will try to predict what are the conditions that can contribute to high severity accidents which may cause loss of life or loss of property .WSOT can use the model to take precaution to minimize the loss of property and life.

- ✚ Reducing the insurance cost and Preventing fatalities

- **Stakeholders:**

- ✚ Government Officials

- ✚ Emergency Responders (911 dispatchers)

- ✚ Common People

- ✚ Insurance Companies

Data understanding

We chose the public data from open source available with labeled columns and attributes and observations data to help us do our analysis better.

Sample Data Below

Link to the Data

<https://s3.us.cloud-object-storage.appdomain.cloud/cf-courses->

<data/CognitiveClass/DPO701EN/version-2/Data-Collisions.csv>

[https://www.seattle.gov/Documents/Departments/S DOT/GIS/Collisions OD.pdf](https://www.seattle.gov/Documents/Departments/S DOT/GIS/Collisions_OD.pdf)

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P
SEVERITYCODE	X	Y	OBJECTID	INKEY	COLDTKEY	REPORTNO	STATUS	ADDRTYPE	INTKEY	LOCATION	EXCEPTS	EXCEPTS	SEVERITYCODE	SEVERITYCODE	COLLISION
2	-122.323	47.70314	1	1307	1307	3502005	Matched	Intersectio	37475	5TH AVE N			2	Injury Coll Angles	
1	-122.347	47.64717	2	52200	52200	2607959	Matched	Block		AURORA BR BETWEEN RAYE ST			1	Property [Sideswip	
1	-122.335	47.60787	3	26700	26700	1482393	Matched	Block		4TH AVE BETWEEN SENECA ST /			1	Property [Parked C	
1	-122.335	47.6048	4	1144	1144	3503937	Matched	Block		2ND AVE E			1	Property [Other	
2	-122.306	47.54574	5	17700	17700	1807429	Matched	Intersectio	34387	SWIFT AVE S AND SWIFT AV OFI			2	Injury Coll Angles	
4	-122.388	47.68858	6	328840	328840	5010477	Matched	Intersectio	36074	34TH AVE			4	Property [Angles	

The data consists of 40 independent variables and 221738 rows. The dependent variable, “SEVERITYCODE”, contains numbers that correspond to different levels of severity caused by an accident from 0 to 4.

Severity codes are as follows:

0: Unknown

1: Property Damage

2: Injury

2b: Serious Injury

3: Fatality

Data Preparation

- Public Traffic data for Seattle city USA is available from the Open source (Link Mentioned above)
- After the data has been extracted, keeping the columns required in the data frame.
- Excluding the rows with null values.
- Transform the data type for analysis.
 - Load the Data to data frame.

Original Size of the Data frame

```
[221738 rows x 40 columns]>  
[5]: #Saving the size of original DF  
print("Size of DF :"+str(df.shape[0])+'x'+str(df.shape[1]))  
Size of DF :221738x40
```

Missing Values from the Data frame

```
#!/: #Check where there are NaNs in the dataframe  
print(df.isnull().sum(axis=0))  
  
X                7478  
Y                7478  
OBJECTID         0  
INCKEY           0  
COLDEKEY         0  
REPORTNO         0  
STATUS           0  
ADORTYPE        3714  
INTKEY          149711  
LOCATION          4593  
EXCEPTSNCODE  120403  
EXCEPTSNDESC  209953  
SEVERITYCODE     1  
SEVERITYDESC     0  
COLLISIONTYPE   26451  
PERSONCOUNT    0  
PEDCOUNT       0  
PEDCYLCOUNT     0  
VEHCOUNT       0  
INJURIES        0  
SERIOUSINJURIES 0  
FATALITIES      0  
INCDATE         0  
INCDTTM         0  
JUNCTIONTYPE    11979  
SDOT_COLCODE    1  
SDOT_COLDESC    1  
INATTENTIONIND  191550  
UNDERINFL       26431  
WEATHER          26641  
ROADCOND        26560  
LIGHTCOND       26730  
PEDROWNOTGRINT  216543  
SDOTCOLNUM      94533  
SPEEDING        211802  
ST_COLCODE      9413  
ST_COLDESC      26451  
SEGLANEKEY      0  
CROSSWALKKEY    0  
HITPARKEDCAR    0  
dtype: int64
```

Preparing the Data after doing the Data Cleansing

```
161: #Data Cleansing  
#find missing WEATHERCOND  
todrop1 = df["WEATHER"] == "Unknown"  
noweatherinfo = todrop1.values.sum()  
  
#find missing ROADCOND  
todrop2 = df["ROADCOND"] == "Unknown"  
noroadcondinfo = todrop2.values.sum()  
  
#find missing LIGHTCOND  
todrop3 = df["LIGHTCOND"] == "Unknown"  
nolightinfo = todrop3.values.sum()  
  
#find missing LIGHTCOND  
todrop4 = df["SEVERITYCODE"] == "2b"  
noseveritycode = todrop4.values.sum()  
  
#Collate these and remove  
df["TODROP"] = 0  
count_noweather = 0  
count_noinfo = 0  
for i in range(0, len(todrop1)):  
    if todrop1[i] == True or todrop2[i] == True or todrop3[i] == True or todrop4[i] == True:  
        df["TODROP"][i] = 1  
  
print("There are "+str(noweatherinfo)+" accidents with no weather information.")  
print("There are "+str(noroadcondinfo)+" accidents with no road condition information.")  
print("There are "+str(nolightinfo)+" accidents with no information about light conditions.")  
print("There are "+str(df["TODROP"].values.sum())+" accidents without one or more of the above.\n Deleting now...")  
  
#Delete the temporary column "TODROP" and re-index the data  
shape0 = df.shape  
todrop = df["TODROP"] == 1  
df.drop(df.index[todrop], inplace=True)
```

```
[7]: #Take columns with mixed boolean data types ([1, "Y", True],[0, "N", False] etc) and cast them as numerical variables
#SPEEDING, INATTENTIONIND, UNDERINFL, PEDROWNOTGRNT, HITPARKEDCAR
df["SPEEDING"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace("Y", 1, inplace=True)

df["INATTENTIONIND"].replace(np.nan, 0, inplace=True)
df["INATTENTIONIND"].replace("Y", 1, inplace=True)

df["INATTENTIONIND"].replace(np.nan, 0, inplace=True)
df["INATTENTIONIND"].replace("Y", 1, inplace=True)

df["UNDERINFL"].replace(np.nan, 0, inplace=True)
df["UNDERINFL"].replace('N', 0, inplace=True)
df["UNDERINFL"].replace('0', 0, inplace=True)
df["UNDERINFL"].replace('1', 1, inplace=True)
df["UNDERINFL"].replace("Y", 1, inplace=True)

df["PEDROWNOTGRNT"].replace(np.nan, 0, inplace=True)
df["PEDROWNOTGRNT"].replace("Y", 1, inplace=True)

df["HITPARKEDCAR"].replace("N", 0, inplace=True)
df["HITPARKEDCAR"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace("Y", 1, inplace=True)
df["HITPARKEDCAR"].replace("Y", 1, inplace=True)
```

Data Cleansing-Dropping the unwanted columns.

```
[13]: if 'OBJECTID' in df:
      del df["OBJECTID"]
      if 'COLDKEY' in df:
          del df["COLDKEY"]
      if 'REPORTNO' in df:
          del df["REPORTNO"]
      if 'STATUS' in df:
          del df["STATUS"]
      if 'EXCEPTSNCODE' in df:
          del df["EXCEPTSNCODE"]
      if 'EXCEPTSNDESC' in df:
          del df["EXCEPTSNDESC"]
      if 'SDOTCOLNUM' in df:
          del df["SDOTCOLNUM"]
      if 'STCOLCODE' in df:
          del df["STCOLCODE"]
      if 'INTKEY' in df:
          del df["INTKEY"]
      if 'LOCATION' in df:
          del df["LOCATION"]
      if 'SEVERITYDESC' in df:
          del df["SEVERITYDESC"]
```

Casting the columns to right datatype (numerical variables) for calculations.


```
[17]: #Take columns with mixed boolean data types ([1, "Y", True],[0, "N", False] etc) and cast them as numerical variables
#SPEEDING, INATTENTIONIND, UNDERINFL, PEDROWNOTGRNT, HITPARKEDCAR
df["SPEEDING"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace("Y", 1, inplace=True)

df["INATTENTIONIND"].replace(np.nan, 0, inplace=True)
df["INATTENTIONIND"].replace("Y", 1, inplace=True)

df["INATTENTIONIND"].replace(np.nan, 0, inplace=True)
df["INATTENTIONIND"].replace("Y", 1, inplace=True)

df["UNDERINFL"].replace(np.nan, 0, inplace=True)
df["UNDERINFL"].replace('N', 0, inplace=True)
df["UNDERINFL"].replace('0', 0, inplace=True)
df["UNDERINFL"].replace('1', 1, inplace=True)
df["UNDERINFL"].replace("Y", 1, inplace=True)

df["PEDROWNOTGRNT"].replace(np.nan, 0, inplace=True)
df["PEDROWNOTGRNT"].replace("Y", 1, inplace=True)

df["HITPARKEDCAR"].replace("N", 0, inplace=True)
df["HITPARKEDCAR"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace(np.nan, 0, inplace=True)
df["SPEEDING"].replace("Y", 1, inplace=True)
df["HITPARKEDCAR"].replace("Y", 1, inplace=True)
```

```
[21]: #Descriptive Stats
descriptive_stats= df.describe(include="all")
```

```
[22]: print(df.SEVERITYCODE.value_counts())
```

```
1    95913
2    49569
3      240
0         2
Name: SEVERITYCODE, dtype: int64
```


Count Based on the Road Condition

```
[25]: print(df.ROADCOND.value_counts())
```

```
Dry          105934
Wet           38373
Ice            687
Snow/Slush     629
Standing Water   50
Sand/Mud/Dirt   30
Oil             21
Name: ROADCOND, dtype: int64
```

```
[27]: df.ROADCOND.value_counts().plot.bar(figsize=(8,3))
```

```
[27]: <AxesSubplot:>
```



Exploratory Data Analysis


We will run a value count on road ('ROADCOND') and weather condition ('WEATHER') to get ideas of the different road and weather conditions. We will also check the value count on light condition ('LIGHTCOND'), to see the breakdowns of accidents occurring during the different light conditions. The results will then be used for data modeling.

```
[22]: print(df.SEVERITYCODE.value_counts())
```

```
1    95913
2    49569
3     240
0         2
Name: SEVERITYCODE, dtype: int64
```

```
[23]: df.SEVERITYCODE.value_counts().plot.bar(figsize=(8,3))
```

```
[23]: <AxesSubplot:>
```




```
[24]: print(df.ROADCOND.value_counts())
```

```
Dry          105934
Wet           38373
Ice            687
Snow/Slush    629
Standing Water  50
Sand/Mud/Dirt  30
Oil            21
Name: ROADCOND, dtype: int64
```


```
[25]: df.ROADCOND.value_counts().plot.bar(figsize=(8,3))
```

```
[25]: <AxesSubplot:>
```




```
[25]: df.WEATHER.value_counts().plot.bar(figsize=(8,3))
```

```
[25]: <AxesSubplot:>
```



```
[24]: df.LIGHTCOND.value_counts().plot.bar(figsize=(8,3))
```

```
[24]: <AxesSubplot:>
```




We built the catplot with seaborn library below to check how the severity is impacted by various attributes.

- Number of pedestrians involved in the collision (PEDCOUNT) and Severity
- Severity 2 (Injury) accidents has happened to Pedestrians, compared to Sev 1 accidents

```
28]: # The number of pedestrians involved in the collision.
import seaborn as sns
sns.catplot(x="SEVERITYCODE", y="PEDCOUNT", data=df, kind="bar")
```

```
28]: <seaborn.axisgrid.FacetGrid at 0x7f2c1bf59198>
```




- Number of people involved in the collision (PERSONCOUNT) and Severity

The data shows that the severity of the accident is high with person count.


```
[27]: # The number of pedestrians involved in the collision.  
import seaborn as sns  
sns.catplot(x="SEVERITYCODE", y="PERSONCOUNT", data=df, kind="bar")
```

```
[27]: <seaborn.axisgrid.FacetGrid at 0x7fa4e90fbda0>
```



- The below data shows accidents are happening more near the blocks and less at intersection. Severity 2 is almost same on Block and on Intersection

```
]# Count plot for Address type :  
sns.countplot(x='ADDRTYPE', data=df, hue='SEVERITYCODE')  
  
]: <AxesSubplot:xlabel='ADDRTYPE', ylabel='count'>
```

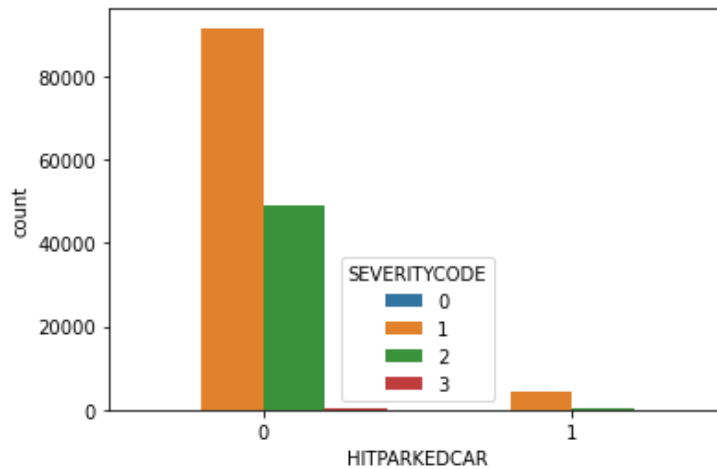


- Number of cases of accidents when car was parked



```
[32]: sns.countplot(x='HITPARKEDCAR', data=df, hue='SEVERITYCODE')
```

```
[32]: <AxesSubplot:xlabel='HITPARKEDCAR', ylabel='count'>
```




Looking at the Speeding cases, there were 9381 speeding cases.

HITPARKEDCAR

```
[51]: sns.countplot(x='SPEEDING', data=df, hue='SEVERITYCODE')
```

```
[51]: <AxesSubplot:xlabel='SPEEDING', ylabel='count'>
```




- Figure shows the high number of accidents when people were under influence

```
[28]: # The number of pedestrians involved in the collision.
```

```
import seaborn as sns
```

```
sns.catplot(x="SEVERITYCODE", y="UNDERINFL", data=df, kind="bar")
```

```
[28]: <seaborn.axisgrid.FacetGrid at 0x7f5eadc8fe10>
```




```
[28]: # The number of pedestrians involved in the collision.
```

```
import seaborn as sns
```

```
sns.catplot(x="SEVERITYCODE", y="UNDERINFL", data=df, kind="bar")
```

```
[28]: <seaborn.axisgrid.FacetGrid at 0x7f5eadc8fe10>
```



We further created an Incident dataset. Following features were used: "Incident Date", "Incident Time", "CollisionDescription", "Weather", "Road condition", "Light Condition" indicating the Incident Detail.

```

26]: #Daily accidents based on the Year
date['year'] = df.INCDATE.dt.year
date['month'] = df.INCDATE.dt.month
date['weekday'] = df.INCDATE.dt.weekday
high_sev = date[date['SEVERITYCODE']==1]


season = date[['INCDATE', 'SEVERITYCODE']].groupby('INCDATE').count()
season['rolling'] = season.SEVERITYCODE.rolling(window=30).mean()
season['SEVERITYCODE'].plot(figsize=(20,10), marker='o', markersize=2, linewidth=0.5, label='Daily trend')
season['rolling'].plot(color='r', linewidth=2, label='rolling mean')
plt.title('Daily accidents, years: 2013, 2014, 2015,2016,2017,2018,2019 and 2020', size=18)
plt.xlabel('Day', size=14)
plt.ylabel('Accidents', size=14)

t0 = dt.datetime.strptime('2012-12-15', '%Y-%m-%d')
t1 = dt.datetime.strptime('2020-10-15', '%Y-%m-%d')

plt.xlim(t0,t1)
plt.legend()


```

The Below graph shows the daily accidents from 2013 till 2020 with high around end of 2016 and beginning of 2017 and the number of accidents low in 2020.



Plotting the accidents Per year, confirms the same trend as above

```
[33]: #By year
yearly = data[['year', 'SEVERITYCODE']].groupby('year').count()
yearly['SEVERITYCODE'].plot.line(figsize=(10,4), marker='o', linewidth=0.5, color='orange', label='Total')
plt.title('Accidents per year')
plt.xticks(range(2012,2020))
plt.xlim(2013,2020)
plt.ylabel('Accidents')
plt.legend()
plt.show()
```



Modeling

- Using Numpy, Scrlar, Linear regression on the clean transformed data (Shown Above)
- After importing necessary packages and splitting preprocessed data into test and train sets, for each machine learning model, we will build and evaluate the model with the techniques as follow:
- Dataframe with features below created
[“ADDRTYPE”, “COLLISIONTYPE”, “JUNCTIONTYPE”, “WEA

THER", "ROADCOND", "LIGHTCOND", "UNDERINFL", "HIT
PARKEDCAR"]]


```
[66]: X = df_balanced.iloc[:,1:]

#Encoding Categorical Features - Training Dataset
X = pd.get_dummies(data=X, columns=['ADDRTYPE', 'COLLISIONTYPE', 'JUNCTIONTYPE', 'WEATHER',
                                     'ROADCOND', 'LIGHTCOND', 'UNDERINFL', 'HITPARKEDCAR'])

Y = df_balanced[['SEVERITYCODE']]
print(X.info())
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, random_state=0)
xtrain, xval, ytrain, yval = train_test_split(X_train, Y_train, test_size=0.25)
print(Y_train.info())
print(Y_test.info())
```

Heatmap

```
[53]: sns.heatmap(df_corr, annot=True, cmap='coolwarm')
plt.savefig('corr.png')
```



Machine Learning Models

- Github as a repository and running Jupyter Notebook are used to process data and build Machine Learning models

- Python and its popular packages such as Pandas, NumPy and Sklearn is used for determining the accuracy.
- The dataset x and y are constructed. After normalization they are split into x_train, y_train, x_test and y_test using train_test_split. 75% of the data is used for training and 25% is used for testing as below

```
[41]: X = df_balanced.iloc[:,1:]

#Encoding Categorical Features - Training Dataset
X = pd.get_dummies(data=X, columns=['ADDRTYPE', 'COLLISIONTYPE', 'JUNCTIONTYPE', 'WEATHER',
                                   'ROADCOND', 'LIGHTCOND', 'UNDERINFL', 'HITPARKEDCAR'])

Y = df_balanced[['SEVERITYCODE']]
print(X.info())
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, random_state=0)
xtrain, xval, ytrain, yval = train_test_split(X_train, Y_train, test_size=0.25)
print(Y_train.info())
print(Y_test.info())
```

- K Nearest Neighbour (KNN)

```
Best Hyperparameter KNN : {'n_neighbors': 6, 'p': 1}
[[ 0  1  0  0]
 [ 0 22073 1991 0]
 [ 0 11024 1280 0]
 [ 0  57  5  0]]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/_classification.py:136: UndefinedMetricWarning: Precision and Recall are ill-defined and being set to 0.0 in labels with no predicted samples
  'precision', 'predicted', average, warn_for)
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.67	0.92	0.77	24064
2	0.39	0.10	0.16	12304
3	0.00	0.00	0.00	62
micro avg	0.64	0.64	0.64	36431
macro avg	0.26	0.26	0.23	36431
weighted avg	0.57	0.64	0.57	36431

```

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/model_selection/_split.py:2053: FutureWarning: You should use 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
  warnings.warn(CV_WARNING, FutureWarning)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/model_selection/_split.py:652: Warning: The least populated class has a minimum number of samples with weight 0.
  % (min_groups, self.n_splits)), Warning)
Best Hyperparameter KNN : {'n_neighbors': 6, 'p': 1}
[[ 0  1  0  0]
 [ 0 22073 1991  0]
 [ 0 11024 1280  0]
 [ 0  57  5  0]]

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.67	0.92	0.77	24064
2	0.39	0.10	0.16	12304
3	0.00	0.00	0.00	62
micro avg	0.64	0.64	0.64	36431
macro avg	0.26	0.26	0.23	36431
weighted avg	0.57	0.64	0.57	36431

```

0.6410200104306771

```

• LINEAR REGRESSION

```

"this warning.", FutureWarning)
['1' '1' '1' ... '1' '1' '1']

[[ 0  1  0  0]
 [ 0 24064 0  0]
 [ 0 12304 0  0]
 [ 0  62  0  0]]

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.66	1.00	0.80	24064
2	0.00	0.00	0.00	12304
3	0.00	0.00	0.00	62
micro avg	0.66	0.66	0.66	36431
macro avg	0.17	0.25	0.20	36431
weighted avg	0.44	0.66	0.53	36431

```

0.6605363563997694

```

• DECISION TREE

We changed our dataset and started again to find the accuracy and which model fits the best.

Dataframe: data clean

```
memory usage: 15.0+ MB

[68]: data_clean = pd.concat([data_clean.drop(['WEATHER', 'ROADCOND', 'LIGHTCOND'], axis=1),
                             pd.get_dummies(data_clean['ROADCOND']),
                             pd.get_dummies(data_clean['LIGHTCOND']),
                             pd.get_dummies(data_clean['WEATHER'])], axis=1)

[69]: data_clean = data_clean.sample(frac=1).reset_index(drop=True)

[70]: from sklearn import preprocessing
x = data_clean.drop(['SEVERITYCODE'], axis=1)
y = data_clean[['SEVERITYCODE']]
data_clean_scaled = preprocessing.StandardScaler().fit(x).transform(x)
data_clean_scaled[0:3]
```

• DECISION TREE

```
[72]: #DECISION TREE ON SECOND SET OF DATA
from sklearn.tree import DecisionTreeClassifier
dTreeModel = DecisionTreeClassifier(criterion='entropy', max_depth=5)
dTreeModel.fit(x_train, y_train)
dTreeModel

[72]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=5,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')

[73]: yHat = dTreeModel.predict(x_test)

[74]: print(classification_report(y_test, yHat))

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision is ill-defined for label 0 with no predicted samples.
'precision', 'predicted', average, warn_for)
      precision    recall  f1-score   support

      0         0.00      0.00      0.00         1
      1         0.71      0.98      0.83       19147
      2         0.86      0.24      0.37        9945
      3         0.00      0.00      0.00         52

 micro avg       0.73      0.73      0.73       29145
 macro avg       0.39      0.31      0.30       29145
 weighted avg    0.76      0.73      0.67       29145

[75]: #Accuracy
acc = accuracy_score(Y_test,yHat)
print(acc,'\n')
accDict['RFT'] = acc

0.6296791902556185
```

```
0]: yHat = dTreeModel.predict(x_test)

1]: print(classification_report(y_test, yHat))

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision is ill-defined for label 0 with no predicted samples.
'precision', 'predicted', average, warn_for)
      precision    recall  f1-score   support

      1         0.72      0.95      0.82       19188
      2         0.75      0.29      0.42        9910
      3         0.00      0.00      0.00         47

 micro avg       0.73      0.73      0.73       29145
 macro avg       0.49      0.41      0.41       29145
 weighted avg    0.73      0.73      0.68       29145

2]: #Accuracy
acc = accuracy_score(Y_test,yHat)
print(acc,'\n')
accDict['RFT'] = acc

0.619488763081146
```

• LINEAR REGRESSION

```

nged to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning: A column-vector
assed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:460: FutureWarning: Default multi_class
e changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.
"this warning.", FutureWarning)
['1' '1' '1' ... '1' '1' '1']

[[18632  556    0]
 [ 7430 2480    0]
 [   16   31    0]]

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precisi
-score are ill-defined and being set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)

```

	precision	recall	f1-score	support
1	0.71	0.97	0.82	19188
2	0.81	0.25	0.38	9910
3	0.00	0.00	0.00	47
micro avg	0.72	0.72	0.72	29145
macro avg	0.51	0.41	0.40	29145
weighted avg	0.75	0.72	0.67	29145

```

0.7243781094527363

```

• RANDOM FOREST

```

: yHat = rfcModel.predict(x_test)

: print(classification_report(y_test, yHat))

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precisi
samples.
'precision', 'predicted', average, warn_for)

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.72	0.96	0.83	19147
2	0.79	0.28	0.42	9945
3	0.00	0.00	0.00	52
micro avg	0.73	0.73	0.73	29145
macro avg	0.38	0.31	0.31	29145
weighted avg	0.74	0.73	0.69	29145

```

: #Accuracy
acc = accuracy_score(yHat,y_test)
print(acc,'\n')
accDict = {}
accDict['RF1'] = acc

0.7301423914908217

```

- KNN

```
[87]: #KNN ON SECOND SET OF DATA
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 8)
classifier.fit(x_train, y_train)

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/ipykernel_launcher.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change
), for example using ravel().
This is separate from the ipykernel package so we can avoid doing imports until
[87]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=8, p=2,
weights='uniform')
[88]: y_pred = classifier.predict(x_test)
[89]: #Accuracy
acc = accuracy_score(y_pred,y_test)
print(acc,'\n')
accDict['KNN'] = acc
0.7234517069823297
```

CONCLUSION

We showed the impact of weather, road condition, road intersection, under influence on the two classes property damage (class 1) or injury (class 2) with graphs.

Except DecisionTree the other three models have a close accuracy of 72-73% which means that the model has trained well but we can have more accuracy with additional features and by having more data for severity code and missing description.

KNN, Logistic Regression, Random forest have an accuracy of 72-74% ,this is possibly the similarity of the features for both types of accidents(1 and 2)

Though with 73% of accuracy from Random Forest we can say that the model has trained well and performs well on the testing as well as the trained data.